

# Unsupervised Traffic Accident Detection in First-Person Videos

Yu Yao<sup>1\*</sup>, Mingze Xu<sup>2\*</sup>, Yuchen Wang<sup>2</sup>, David J. Crandall<sup>2</sup>, Ella M. Atkins<sup>1</sup>

**Abstract**—Recognizing abnormal events such as traffic violations and accidents in natural driving scenes is essential for successful autonomous driving and advanced driver assistance systems. However, most work on video anomaly detection suffers from two crucial drawbacks. First, they assume cameras are fixed and videos have static backgrounds, which is reasonable for surveillance applications but not for vehicle-mounted cameras. Second, they pose the problem as one-class classification, relying on arduously hand-labeled training datasets that limit recognition to anomaly categories that have been explicitly trained. This paper proposes an unsupervised approach for traffic accident detection in first-person (dashboard-mounted camera) videos. Our major novelty is to detect anomalies by predicting the future locations of traffic participants and then monitoring the prediction accuracy and consistency metrics with three different strategies. We evaluate our approach using a new dataset of diverse traffic accidents, AnAn Accident Detection (A3D), as well as another publicly-available dataset. Experimental results show that our approach outperforms the state-of-the-art. Code and the dataset developed in this work are available at: <https://github.com/MoonBlvd/tad-IROS2019>

## I. INTRODUCTION

Autonomous driving has the potential to transform the world as we know it, revolutionizing transportation by making it faster, safer, cheaper, and less labor intensive. A key challenge is building autonomous systems that can accurately perceive and safely react to the huge diversity in situations that are encountered on real-world roadways. Driving situations obey a long-tailed distribution, such that a very small number of common situations makes up the vast majority of what a driver encounters, and a virtually infinite number of rare scenarios — animals running into the roadway, cars driving on the wrong side of the street, etc. — makes up the rest. While each of these individual scenarios is rare, they can and do happen. In fact, the chances that *one* of them will occur on any given day are actually quite high.

Existing work in computer vision has applied deep learning-based visual classification to detect action starts and their associated categories [1] in the video collected by dashboard-mounted cameras [2]. The long-tailed distribution of driving events means that unusual events may occur so infrequently that it may be impossible to collect training data for them, or to even anticipate that they might occur [3]. In fact, some studies indicate that driverless cars would need to be tested for *billions* of miles before enough of these rare situations occur to even accurately measure system safety [4],

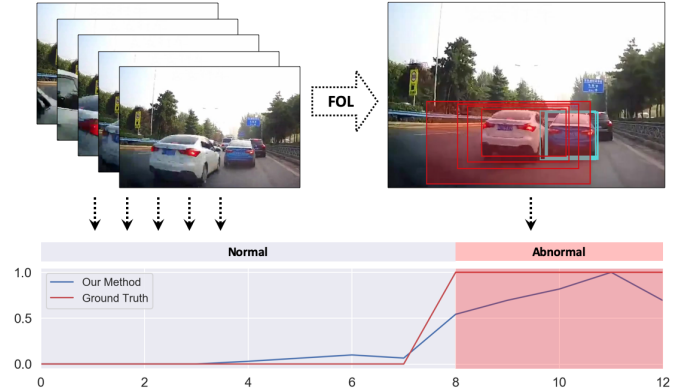


Fig. 1: Overview of our proposed approach. For each time  $t$ , we monitor the accuracy and consistency of all traffic participants' predicted bounding boxes from previous frames and calculate the scene's anomaly score.

much less to collect sufficient training data to make them work well.

An alternative approach is to avoid modeling all possible driving scenarios, but instead to train models that recognize “normal,” safe roadway conditions, and then signal an anomaly when events that do not fit the model are observed. Unlike the fully-supervised classification-based work, this unsupervised approach would not be able to identify exactly which anomaly has occurred, but it may still provide sufficient information for the driving system to recognize an unsafe situation and take evasive action. This paper proposes a novel approach that learns a deep neural network model to predict the future locations of objects such as cars, bikes, pedestrians, etc., in the field of view of a dashboard-mounted camera on a moving ego-vehicle. These models can be easily learned from massive collections of dashboard-mounted video of normal driving, and no manual labeling is required. We then compare predicted object locations to the actual locations observed in the next few video frames. We hypothesize that anomalous roadway events can be detected by looking for major deviations between the predicted and actual locations, because unexpected roadway events (such as cars striking other objects) result in sudden unexpected changes in an object's speed or position.

Perhaps the closest related work to ours is Liu *et al.* [3], who also detect anomalous events in video. Their technique tries to predict entire future RGB frames and then looks for deviations between those and observed RGB frames. But while their approach can work well for static cameras, accurately predicting whole frames is extremely difficult when cameras are rapidly moving, as in the driving scenario.

\*The first two authors contributed equally.

<sup>1</sup>Robotics Institute, University of Michigan, Ann Arbor, MI 48109, USA. {brianyao, ematkins}@umich.edu

<sup>2</sup>School of Informatics, Computing, and Engineering, Indiana University, Bloomington, IN 47408, USA. {mx6, wang617, djcran}@iu.edu

We side-step this difficult problem by detecting objects and predicting their trajectories, as opposed to trying to predict whole frames. To model the moving camera, we explicitly predict the future odometry of the ego-vehicle; this also allows us to detect significant deviations of the predicted and real ego-motion, which can be used to classify if the ego-vehicle is involved in the accident or is just an observer. We evaluate our technique in extensive experiments on three datasets, including a new labeled dataset of some 1,500 video traffic accidents from dashboard cameras that we collected from YouTube. We find that our method significantly outperforms a number of baselines, including the published state-of-the-art in anomaly detection.

## II. RELATED WORK

**Trajectory Prediction.** Extensive research has investigated trajectory prediction, often posed as a sequence-to-sequence generation problem. Alahi *et al.* [5] introduce a Social-LSTM for pedestrian trajectories and their interactions. The proposed social pooling method is further improved by Gupta *et al.* [6] to capture global context in a Generative Adversarial Network (GAN). Social pooling is also applied to vehicle trajectory prediction in Deo *et al.* [7] with multi-modal maneuver conditions. Other work [8], [9] captures scene context information using attention mechanisms to assist trajectory prediction. Lee *et al.* [10] incorporate Recurrent Neural Networks (RNNs) with conditional variational autoencoders (CVAEs) to generate multimodal predictions and choose the best by ranking scores.

While the above methods are designed for third-person views from static cameras, recent work has considered vision in first-person (egocentric) videos that capture the natural field of view of the person or agent (e.g., vehicle) wearing the camera to study the camera wearer's actions [11], [12], trajectories [13], interactions [14], [15], etc. Bhattacharyya *et al.* [16] predict future locations of pedestrians from vehicle-mounted cameras, modeling observation uncertainties with a Bayesian LSTM network. Yagi *et al.* [17] incorporate different kinds of cues into a convolution-deconvolution (Conv1D) network to predict pedestrians' future locations. Yao *et al.* [18] extend this work to autonomous driving scenarios by proposing a multi-stream RNN Encoder-Decoder (RNN-ED) architecture with both past vehicle locations and image features as inputs for anticipating vehicle locations.

**Video Anomaly Detection.** Video anomaly detection has received considerable attention in computer vision and robotics [19]. Previous work mainly focuses on video surveillance scenarios typically using an unsupervised learning method on the reconstruction of normal training data. For example, Hasan *et al.* [20] propose a 3D convolutional Auto-Encoder (Conv-AE) to model non-anomalous frames. To take advantage of temporal information, [21], [22] use a Convolutional LSTM Auto-Encoder (ConvLSTM-AE) to capture regular visual and motion patterns simultaneously. Luo *et al.* [23] propose a special framework of sRNN, called temporally-coherent sparse coding (TSC), to preserve

the similarities between frames within normal and abnormal events. Liu *et al.* [3] detect anomalies by looking for differences between a predicted future frame and the actual frame. However, in dynamic autonomous driving scenarios, it is hard to reconstruct either the current or future RGB frames due to the ego-car's intense motion. It is even harder to detect abnormal events. This paper proposes detecting accidents on roads by using the difference between predicted and actual trajectories of other vehicles. Our method not only eliminates the computational cost of reconstructing full RGB frames, but also localizes potential anomaly participants.

Prior work has also detected anomalies such as moving violations and car collisions on roads. Chan *et al.* [2] introduce a dataset of crowd-sourced dashcam videos and a dynamic-spatial-attention RNN model for accident detection. Herzig *et al.* [24] propose a Spatio-Temporal Action Graph (STAG) network to model the latent graph structure of spatial and temporal relations between objects. These methods are based on supervised learning that requires arduous human annotations and makes the unrealistic assumption that all abnormal patterns have been observed in the training data. This paper considers the challenging but practical problem of predicting accidents with unsupervised learning. To evaluate our approach, we introduce a new dataset with traffic accidents involving objects such as cars and pedestrians.

## III. UNSUPERVISED TRAFFIC ACCIDENT DETECTION IN FIRST-PERSON VIDEOS

Autonomous vehicles must monitor the roadway ahead for signs of unexpected activity that may require evasive action. A natural way to detect these anomalies is to look for unexpected or rare movements in the first-person perspective of a front-facing, dashboard-mounted camera on a moving ego-vehicle. Prior work [3] proposes monitoring for unexpected scenarios by using past video frames to predict the current video frame, and then comparing it to the observed frame and looking for major differences. However, this does not work well for moving cameras on vehicles, where the perceived optical motion in the frame is induced by both moving objects and camera ego-motion. More importantly, anomaly detection systems do *not* need to accurately predict all information in the frame, since anomalies are unlikely to involve peripheral objects such as houses or billboards by the roadside. This paper thus assumes that an anomaly may exist if an object's real-world observed trajectory deviates from the predicted trajectory. For example, when a vehicle should move through an intersection but instead suddenly stops, a collision may have occurred.

Following Liu *et al.* [3], our model is trained with a large-scale dataset of normal, non-anomalous driving videos. This allows the model to learn normal patterns of object and ego motions, then recognize deviations without the need to explicitly train the model with examples of every possible anomaly. This video data is easy to obtain and does not require hand labeling. Considering the influence of ego-motion on perceived object location, we incorporate a future ego-motion prediction module [18] as an additional input. At



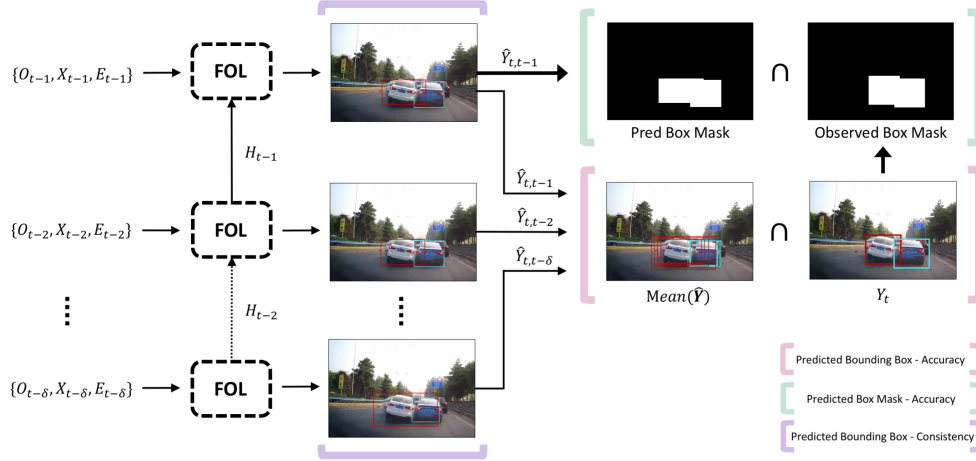


Fig. 3: Overview of our unsupervised traffic accident detection methods. The three brackets correspond to: (1) Predicted bounding box accuracy method (pink); (2) Predicted box mask accuracy method (green); (3) Predicted bounding box consistency method (purple). All methods use multiple previous FOL outputs to compute anomaly scores.

curacy and consistency of objects' future locations. The key idea is that object trajectories and locations in non-anomalous events can be precisely predicted, while deviations from predicted behaviors suggest an anomaly.

1) *Predicted Bounding Boxes - Accuracy*: One simple method for recognizing abnormal events is to directly measure the similarity between predicted object bounding boxes and their corresponding observations. The FOL model predicts bounding boxes of the next  $\delta$  future frames, i.e., at each time  $t$  each object has  $\delta$  bounding boxes predicted from time  $t - \delta$  to  $t - 1$ , respectively. We first average the positions of the  $\delta$  bounding boxes, then compute intersection over union (IoU) between the averaged bounding box and the observed box location, where higher IoU means greater agreement between the two boxes. We average computed IoU values over all observed objects and then compute an aggregate anomaly score  $L_{bbox} \in [0, 1]$ ,

$$L_{bbox} = 1 - \frac{1}{N} \sum_{i=1}^N \text{IoU} \left( \left( \frac{1}{\delta} \sum_{j=1}^{\delta} \hat{Y}_{t,t-j}^i \right), Y_{t_0}^i \right), \quad (1)$$

where  $N$  is the total number of observed objects, and  $\hat{Y}_{t,t-j}^i$  is the predicted bounding box from time  $t - j$  of object  $i$  at time  $t$ . This method relies upon accurate object tracking to match the predicted and observed bounding boxes.

2) *Predicted Box Mask - Accuracy*: Although tracking algorithms such as Deep-SORT [26] offer reasonable accuracy, it is still possible to lose or mis-track objects. We found that inaccurate tracking particularly happens in severe traffic accidents because of the twist and distortion of object appearances. Moreover, severe ego-motion also results in inaccurate tracking due to sudden changes in object locations. This increases the number of false negatives of the metric proposed above, which simply ignores objects that are not successfully tracked in a given frame. To solve this problem, we first convert all areas within the predicted bounding boxes to binary masks, with areas inside the boxes having value 1 and backgrounds having 0, and do the same

with the observed boxes. We then calculate an anomaly score as the IoU between these two binary masks,

$$I^{(u,v)} = \begin{cases} 1, & \text{if pixel } (u,v) \text{ within box } X^i, \forall i, \\ 0, & \text{otherwise,} \end{cases} \quad (2)$$

$$L_{mask} = 1 - \text{IoU}(\hat{I}_{t,t-1}, I_t), \quad (3)$$

where  $I^{(u,v)}$  is pixel  $(u,v)$  on mask  $I$ ,  $X^i$  is the  $i$ -th bounding box,  $\hat{I}_{t,t-1}$  is the predicted mask from time  $t - 1$ , and  $I_t$  is the observed mask at  $t$ . In other words, while the metric in the last section compares bounding boxes on an object-by-object basis, this metric simply compares the bounding boxes of all objects simultaneously. The main idea is that accurate prediction results will still have a relatively large IoU compared to the ground truth observation.

3) *Predicted Bounding Boxes - Consistency*: The above methods rely on accurate detection of objects in consecutive frames to compute anomaly scores. However, the detection of anomaly participants is not always accurate due to changes in appearance and mutual occlusions. We hypothesize that visual and motion features about an anomaly do not only appear once it happens, but usually are accompanied by a salient pre-event. We thus propose another strategy to detect anomalies by computing consistency of future object localization outputs from several previous frames while eliminating the effect of inaccurate detection and tracking.

As discussed in Section III-B.1, our model has  $\delta$  predicted bounding boxes for each object in video frame  $t$ . We compute the standard deviation (STD) between all  $\delta$  predicted bounding boxes to measure their similarity,

$$L_{pred} = \frac{1}{N} \sum_{i=1}^N \max_{\{c^x, c^y, w, h\}} \text{STD}(\hat{Y}_{t,t-j}). \quad (4)$$

We compute the maximum STD over the four components of the bounding boxes since different anomalies may be indicated by different effects on the bounding box, e.g., suddenly stopped cross traffic may only have large STD



TABLE I: Comparison of publicly available datasets for video anomaly detection. \*Surveillance videos. \*\*Egocentric videos (training frames are all normal videos, while some test frame videos contain anomalies.)

Dataset	# videos	# training frames	# testing frames	# anomaly events	typical participants
UCSD Ped1/Ped2* [27]	98	9,350	9,210	77	bike, pedestrian, cart, skateboard
CUHK Avenue* [28]	37	15,328	15,324	47	bike, pedestrian
UCF-Crime* [29]	1,900	1,610 videos	290 videos	1,900	car, pedestrian, animal
ShanghaiTech* [23]	437	274,515	42,883	130	bike, pedestrian
Street Accidents (SA)** [2]	994	82,900	16,500	165	car, truck, bike
<b>A3D**</b>	<b>1,500</b>	<b>79,991 (HEV-I)</b>	<b>128,175</b>	<b>1,500</b>	<b>car, truck, bike, pedestrian, animal</b>

along the horizontal axis. A low STD suggests the object is following normal movement patterns and thus the predictions are stable, while a high standard deviation suggests abnormal motion. For all three methods, we follow [3] to normalize computed anomaly scores for evaluation.

#### IV. EXPERIMENTS

To evaluate our method on realistic traffic scenarios, we introduce a new dataset, **AnAn Accident Detection (A3D)**, of on-road abnormal event videos compiled as 1500 video clips from a YouTube channel [30] of dashboard cameras from different cars in East Asia. Each video contains an abnormal traffic event at different temporal locations. We labeled each video with anomaly start and end times under the consensus of three human annotators. The annotators were instructed to label the anomalies based on common sense, with the start time defined to be the point when the accident is inevitable and the end time the point when all participants recover a normal *moving* condition or fully stop.

We compare our A3D dataset with existing video anomaly detection datasets in Table I. A3D includes a total of 128,175 frames (ranging from 23 to 208 frames) at 10 frames per second and is clustered into 18 types of traffic accidents each labeled with a brief description. A3D includes driving scenarios with different weather conditions (e.g., sunny, rainy, snowy, etc.), places (e.g., urban, countryside, etc.), and participant types (e.g., cars, motorcycles, pedestrians, animals, etc.). In addition to start and end times, each traffic anomaly is labeled with a binary value indicating whether the ego-vehicle is involved, to provide a better understanding of the event. Note that this could especially benefit the first-person vision community. For example, rear-end collisions are the most difficult to detect from traditional anomaly detection methods. About 60% of accidents in the dataset involve the ego-vehicle, and others are observed by moving cars from a third-person perspective.

Since A3D does not contain nominal videos, we use the publicly available Honda Egocentric View Intersection (HEV-I) [18] dataset to train our model. HEV-I was designed for future object localization and consists of 230 on-road videos at intersections in the San Francisco Bay Area. Each video is 10-60 seconds in length. Since HEV-I and A3D were collected in different places with different kinds of cameras, there is no overlap between the training and testing datasets. Following prior work [18], we produce object bounding

boxes using Mask-RCNN [31] pre-trained on the COCO dataset and find tracking IDs using Deep-SORT [26].

##### A. Implementation Details

We implemented our model in PyTorch [32] and performed experiments on a system with an Nvidia Titan Xp Pascal GPU. We use ORB-SLAM 2.0 [33] for ego odometry calculation and compute optical flow using FlowNet 2.0 [34]. In our training data (HEV-I), we used the provided camera intrinsic matrix. We used the same matrix for A3D and SA since these videos are collected from different dash cameras and the parameters are unavailable. We also set the feature count to 12000 to have a better performance. We use a  $5 \times 5$  RoIPool operator to produce the final flattened feature vector  $O_t \in \mathbb{R}^{50}$ . The gated recurrent unit (GRU) [35] is our basic RNN cell. GRU hidden state sizes for future object localization and the ego-motion prediction model were set to 512 and 128, respectively. To learn network parameters, we use the RMSprop [36] optimizer with default parameters, learning rate  $10^{-4}$ , and no weight decay. Our models were optimized in an end-to-end manner, and the training process was terminated after 100 epochs using a batch size of 32. The best model was selected according to its performance in future object localization.

##### B. Evaluation Metrics

For accident detection evaluation, we follow the literature of video anomaly detection [27] and compute frame-level Receiver Operation Characteristic (ROC) curves and Area Under the Curve (AUC). A higher AUC value indicates better performance.

##### C. Video Anomaly Detection Baselines

**K-Nearest Neighbor Distance.** We segment each video into a bag of short video chunks of 16 frames. Each chunk is labeled as either normal or anomalous based on the annotation of the 8-th frame. We then feed each chunk into an I3D [37] network pre-trained on Kinetics dataset, and extract the outputs of the last fully connected layer as its feature representations. All videos in the HEV-I dataset are used as normal data. The normalized distance of each test video chunk to the centroid of its  $K$  nearest normal ( $K$ -NN) video chunks is computed as the anomaly score. We show results of  $K = 1$  and  $K = 5$  in this paper.

**Conv-AE [20].** We reimplement the Conv-AE model for unsupervised video anomaly detection by following [20]. The

TABLE II: Experimental results of FOL (errors are in pixels).

Dataset	Prediction Horizon	FDE	ADE	FIOU
HEV-I (test) [18]	0.5 sec	11.0	6.7	0.85
SA (test) [2]	0.5 sec	21.3	13.5	0.64
A3D	0.5 sec	25.6	16.4	0.63

TABLE III: Experimental results on A3D and SA datasets in terms of AUC.

Methods	A3D	A3D (w/o Ego)	SA [2]
$K$ -NN ( $K = 1$ )	48.0	51.3	48.2
$K$ -NN ( $K = 5$ )	47.8	51.2	48.1
Conv-AE [20]	49.5	49.9	50.4
State-of-the-art [3]	46.1	50.7	50.4
FOL-AvgIoU	49.7	57.0	53.4
FOL-MinIoU	48.4	56.0	52.6
FOL-Mask	54.1	54.9	54.8
FOL-AvgSTD (pred only)	59.3	<b>60.2</b>	<b>55.8</b>
FOL-MaxSTD (pred only)	<b>60.1</b>	59.8	55.6

input images are encoded by 3 convolutional layers and 2 pooling layers, and then decoded by 3 deconvolutional layers and 2 upsampling layers for reconstruction. Anomaly score computation is from [20]. The model is trained on a mixture of the SA (Table I) and the HEV-I dataset for 20 epochs and the best model is selected.

**State-of-the-art [3].** The future frame prediction network with Generative Adversarial Network (GAN) achieved the state-of-the-art results for video anomaly detection. This work detects abnormal events by leveraging the difference between a predicted future frame and its ground truth. To fairly compare with our method, we used the publicly available code by the authors of [3] and finetuned on the same dataset as Conv-AE. Training is terminated after 100,000 iterations and the best model is selected.

#### D. FOL Results

We first show the performance of the pretrained FOL model on HEV-I's validation set and on the other two accident datasets (SA and A3D). Similar to [18], the final displacement error (FDE), average displacement error (ADE), and final IOU (FIOU) are presented in Table II. The FDEs and ADEs on A3D and SA are higher and the FIOUs are lower than HEV-I because these videos were collected from different dash cameras in different scenarios, while all HEV-I videos were collected using the same cameras. The accidents in these videos result in lower FOL prediction accuracy, which is consistent with the assumption of our proposed approach. The overall FOL performance on A3D is slightly worse compare to SA since A3D is a larger dataset with more diverse accident types.

#### E. Accident Detection Results on A3D Dataset

**Quantitative Results.** We evaluated baselines, a state-of-the-art method, and our proposed method on the A3D dataset.

As shown in the first column of Table III, our method outperforms the  $K$ -NN baseline as well as Conv-AE and the state-of-the-art. As a comparative study, we evaluate performance of our future object localization (FOL) methods with the three metrics presented in Section III-B. *FOL-AvgIoU* uses the metrics in Eq. (1), while *FOL-MinIoU* is a variation where we evaluate minimum IoU over all observed objects instead of computing the average, resulting in not only anomaly detection but also anomalous object localization. However, *FOL-MinIoU* can perform worse since it is not robust to outliers such as failed prediction of a normal object, which is more frequent in videos with a large number of objects. *FOL-Mask* uses the metrics in Eq. (3) and significantly outperforms the above two methods. This method does not rely on accurate tracking, so it handles cases including mis-tracked objects. However, it may mis-label a frame as an anomaly if object detection loses some normal objects. Our best methods use the prediction-only metric defined in Eq. (4) which has two variations *FOL-AvgSTD* and *FOL-MaxSTD*. Similar to the IoU based methods, *FOL-MaxSTD* finds the most anomalous object in the frame. By using only prediction, our method is insensitive to unreliable object detection and tracking when an anomaly happens, including the false negatives (in IoU based methods) and the false positives (in Mask based methods) caused by losing objects. However, this method can fail in cases where predicting future locations of an object is difficult, e.g., an object with low resolution, intense ego-motion, or multiple object occlusions due to heavy traffic.

We also evaluated the methods by removing videos where ego cars (A3D w/o Ego in Table III) are involved in anomalies to show how ego motion influences anomaly detection performance. As shown in the second and the third columns of Table III, *FOL-AvgIoU* and *FOL-MinIoU* perform better on videos where ego camera is steady while the other methods are relatively robust to ego-motion. This further shows that it is necessary to reduce dependency on accurate object detection and tracking when anomalies occur.

**Qualitative Results.** Fig. 4 shows two sample results of our best method and the published state-of-the-art on the A3D dataset. For example, in the upper one, predictions of all observed traffic participants are accurate and consistent at the beginning. The ego car is hit at around the 30-th frame by the white car on its left, causing inaccurate and unstable predictions and generating high anomaly scores. After the crash, the ego car stops and the predictions recover, as presented in the last two images. Fig. 6 shows a failure case where our method makes false alarms at the beginning due to inconsistent prediction of the very left car occluded by trees. This is because our model takes all objects into consideration equally rather than focusing on important objects. False negatives show that our method is not able to detect an accident if participants are totally occluded (e.g. the bike) or the motion pattern is accidentally normal from a particular viewpoint (e.g. the middle car).




Fig. 4: Two examples of our best method and a state-of-the-art method on the A3D dataset.

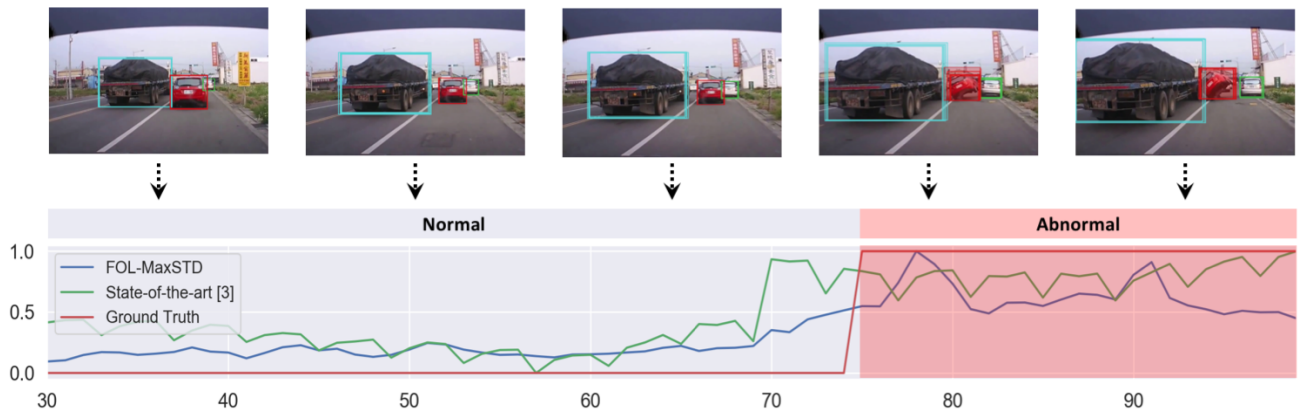


Fig. 5: An example of our best method and a state-of-the-art method on the SA dataset [2].

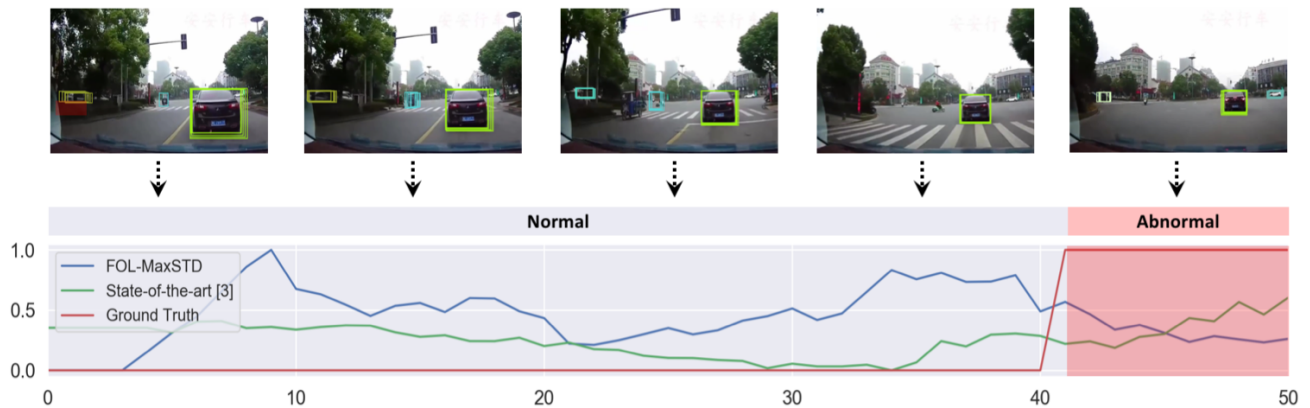


Fig. 6: A failure case of our method on the A3D dataset with false alarms and false negatives.



## F. Results on the SA Dataset

We also compared the performance of our model and baselines on the Street Accident (SA) [2] dataset of on-road accidents in Taiwan. This dataset was collected from dashboard cameras with 720p resolution from the driver's point-of-view. Note that we use SA only for testing, and still train on the HEV-I dataset. We follow prior work [2] and report evaluation results with 165 test videos containing different anomalies. The right-most column in Table III shows the results of different methods on SA. In general, our best method outperforms all baselines and the published state-of-the-art. The SA testing dataset is much smaller than A3D, and we have informally observed that it is biased towards anomalies involving bikes. It also contains videos collected from cyclist head cameras which have irregular camera angles and large vibrations. Fig. 5 shows an example of anomaly detection in the SA dataset.

## V. CONCLUSION

This paper proposed an unsupervised deep learning framework for traffic accident detection from egocentric videos. A key challenge is rapid motion of the ego-car, making visual reconstruction of either current or future RGB frames from regular training data difficult. We predicted traffic participant trajectories as well as their future locations, and utilized anticipation accuracy and consistency as signals that an anomaly may have occurred. We introduced a new dataset consisting of a variety of real-world accidents on roads and also evaluated our method on an existing traffic accident detection dataset. Experiments showed that our model significantly outperforms published baselines.

## VI. ACKNOWLEDGMENTS

This research has been supported by the National Science Foundation under awards CNS 1544844 and CAREER IIS-1253549, and by the IU Office of the Vice Provost for Research, the College of Arts and Sciences, and the School of Informatics, Computing, and Engineering through the Emerging Areas of Research Project "Learning: Brains, Machines, and Children." The views and conclusions contained in this paper are those of the authors and should not be interpreted as representing the official policies, either expressly or implied, of the U.S. Government, or any sponsor.

## REFERENCES

- [1] M. Gao, M. Xu, L. S. Davis, R. Socher, and C. Xiong, "Startnet: Online detection of action start in untrimmed videos," *ICCV*, 2019.
- [2] F.-H. Chan, Y.-T. Chen, Y. Xiang, and M. Sun, "Anticipating accidents in dashcam videos," in *ACCV*, 2016.
- [3] W. Liu, W. Luo, D. Lian, and S. Gao, "Future frame prediction for anomaly detection—a new baseline," in *CVPR*, 2018.
- [4] N. Kalra and S. M. Paddock, "Driving to safety: How many miles of driving would it take to demonstrate autonomous vehicle reliability?" *Transportation Research Part A: Policy and Practice*, 2016.
- [5] A. Alahi, K. Goel, V. Ramanathan, A. Robicquet, L. Fei-Fei, and S. Savarese, "Social LSTM: Human trajectory prediction in crowded spaces," in *CVPR*, 2016.
- [6] A. Gupta, J. Johnson, L. Fei-Fei, S. Savarese, and A. Alahi, "Social GAN: Socially acceptable trajectories with generative adversarial networks," in *CVPR*, 2018.
- [7] N. Deo, A. Rangesh, and M. M. Trivedi, "How would surround vehicles move? a unified framework for maneuver classification and motion prediction," *T-IV*, 2018.
- [8] A. Sadeghian, F. Legros, M. Voisin, R. Vesel, A. Alahi, and S. Savarese, "Car-Net: Clairvoyant attentive recurrent network," in *ECCV*, 2018.
- [9] A. Sadeghian, V. Kosaraju, A. Sadeghian, N. Hirose, and S. Savarese, "Sophie: An attentive gan for predicting paths compliant to social and physical constraints," *arXiv:1806.01482*, 2018.
- [10] N. Lee, W. Choi, P. Vernaza, C. B. Choy, P. H. Torr, and M. Chandraker, "Desire: Distant future prediction in dynamic scenes with interacting agents," in *CVPR*, 2017.
- [11] Y. Li, Z. Ye, and J. M. Rehg, "Delving into egocentric actions," in *CVPR*, 2015.
- [12] M. Ma, H. Fan, and K. M. Kitani, "Going deeper into first-person activity recognition," in *CVPR*, 2016.
- [13] G. Bertasius, A. Chan, and J. Shi, "Egocentric basketball motion planning from a single first-person image," in *CVPR*, 2018.
- [14] C. Fan, J. Lee, M. Xu, K. K. Singh, Y. J. Lee, D. J. Crandall, and M. S. Ryoo, "Identifying first-person camera wearers in third-person videos," *CVPR*, 2017.
- [15] M. Xu, C. Fan, Y. Wang, M. S. Ryoo, and D. J. Crandall, "Joint person segmentation and identification in synchronized first-and third-person videos," *ECCV*, 2018.
- [16] A. Bhattacharyya, M. Fritz, and B. Schiele, "Long-term on-board prediction of people in traffic scenes under uncertainty," in *CVPR*, 2018.
- [17] T. Yagi, K. Mangalam, R. Yonetani, and Y. Sato, "Future person localization in first-person videos," in *CVPR*, 2018.
- [18] Y. Yao, M. Xu, C. Choi, D. J. Crandall, E. M. Atkins, and B. Dariush, "Egocentric vision-based future vehicle localization for intelligent driving assistance systems," *ICRA*, 2019.
- [19] V. Chandola, A. Banerjee, and V. Kumar, "Anomaly detection: A survey," in *ACM CSUR*, 2009.
- [20] M. Hasan, J. Choi, J. Neumann, A. K. Roy-Chowdhury, and L. S. Davis, "Learning temporal regularity in video sequences," in *CVPR*, 2016.
- [21] J. R. Medel and A. Savakis, "Anomaly detection in video using predictive convolutional long short-term memory networks," *arXiv:1612.00390*, 2016.
- [22] Y. S. Chong and Y. H. Tay, "Abnormal event detection in videos using spatiotemporal autoencoder," in *ISNN*, 2017.
- [23] W. Luo, W. Liu, and S. Gao, "A revisit of sparse coding based anomaly detection in stacked rnn framework," in *ICCV*, 2017.
- [24] R. Herzig, E. Levi, H. Xu, E. Brosh, A. Globerson, and T. Darrell, "Classifying collisions with spatio-temporal action graph networks," *arXiv:1812.01233*, 2018.
- [25] M. Xu, M. Gao, Y.-T. Chen, L. S. Davis, and D. J. Crandall, "Temporal recurrent networks for online action detection," *ICCV*, 2018.
- [26] N. Wojke, A. Bewley, and D. Paulus, "Simple online and realtime tracking with a deep association metric," in *ICIP*, 2017.
- [27] W. Li, V. Mahadevan, and N. Vasconcelos, "Anomaly detection and localization in crowded scenes," *TPAMI*, 2014.
- [28] C. Lu, J. Shi, and J. Jia, "Abnormal event detection at 150 fps in matlab," in *ICCV*, 2013.
- [29] W. Sultani, C. Chen, and M. Shah, "Real-world anomaly detection in surveillance videos," in *CVPR*, 2018.
- [30] <https://www.youtube.com/channel/UC-Oa3wml6F3YcptlFwaLgDA/featured>.
- [31] K. He, G. Gkioxari, P. Dollár, and R. Girshick, "Mask R-CNN," in *ICCV*, 2017.
- [32] <http://pytorch.org/>.
- [33] R. Mur-Artal and J. D. Tardós, "Orb-slam2: An open-source slam system for monocular, stereo, and rgb-d cameras," *T-RO*, 2017.
- [34] E. Ilg, N. Mayer, T. Saikia, M. Keuper, A. Dosovitskiy, and T. Brox, "FlowNet 2.0: Evolution of optical flow estimation with deep networks," in *CVPR*, 2017.
- [35] J. Chung, C. Gulcehre, K. Cho, and Y. Bengio, "Gated feedback recurrent neural networks," in *ICML*, 2015.
- [36] G. Hinton, N. Srivastava, and K. Swersky, "Neural networks for machine learning, lecture 6a: Overview of mini-batch gradient descent."
- [37] J. Carreira and A. Zisserman, "Quo vadis, action recognition? a new model and the kinetics dataset," in *CVPR*, 2017.